

Topic 8: K-Means Clustering

GO GREEN. AVOID PRINTING, OR PRINT DOUBLE-SIDED.

8.1 Introduction

Recall that classification can be summarized as assigning a label (class) $y \in \{1, 2, \dots, C\} =: [C]$ to a data point $\mathbf{x} \in \mathbb{R}^D$ based on a collection of training data points $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N \in \mathbb{R}^D$ whose corresponding classes $y_1, y_2, \dots, y_N \in [C]$ are already known. To this end we can use nearest neighbors and other *supervised* learning algorithms.

However, in data science the labels $\{y_i\}_{i=1}^N$ of the training data points $\{\mathbf{x}_i\}_{i=1}^N$ are often unavailable:

- Given a collection of vectorized images $\{\mathbf{x}_i\}_{i=1}^N$, we want to determine which correspond to the same individuals (but we don't know their names).
- Given a collection of vectors $\{\mathbf{x}_i\}_{i=1}^N$ containing information about people's movies ratings (e.g., Netflix or Amazon), we want to determine which people have similar preferences.
- Given a collection of vectors $\{\mathbf{x}_i\}_{i=1}^N$ containing genomic sequences from different organisms in a human gut microbiome sample, determine which sequences correspond to the same species.

Unsupervised learning refers to the tasks when labels $\{y_i\}_{i=1}^N$ are unavailable. Clustering is one of such tasks.

8.2 Clustering

The task of clustering can be summarized as splitting a collection of data points into groups such that the points in each group are *similar*. More precisely, given a collection of data points $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N \in \mathbb{R}^D$, we want to identify a partition $\{C_1, C_2, \dots, C_K\}$ of $[N]$ (called clusters) such that if $i, j \in C_k$, then \mathbf{x}_i and \mathbf{x}_j are *close* to each other (recall that there are several ways to define how *close* two points are, e.g., $\|\mathbf{x}_i - \mathbf{x}_j\|_2$, $\|\mathbf{x}_i - \mathbf{x}_j\|_1$, or $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$).

Recall that norms satisfy the so-called triangle inequality: $\|\mathbf{x}\| + \|\mathbf{y}\| \geq \|\mathbf{x} + \mathbf{y}\|$, which implies that if \mathbf{x} is close to \mathbf{y} , and \mathbf{y} is close to \mathbf{z} , then \mathbf{x} is also close to \mathbf{z} . Using this insight, we can rephrase/adapt our clustering goal in terms of *centers* as follows: given a collection of data points $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N \in \mathbb{R}^D$, we want to identify *centers* $\boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \dots, \boldsymbol{\mu}_K \in \mathbb{R}^D$ that minimize the within-cluster distances:

$$\sum_{k=1}^K \sum_{i \in C_k} \|\boldsymbol{\mu}_k - \mathbf{x}_i\|,$$




Figure 8.1: Initial and final step of Lloyd's algorithm.

where $i \in C_k$ if $\|\mu_k - \mathbf{x}_i\| \leq \|\mu_\ell - \mathbf{x}_i\|$ for every $\ell \in [K]$. Notice that this is a kind of chicken and egg problem: you need to know the clusters $\{C_k\}$ to find the centers $\{\mu_k\}$, and you need to know the centers in order to determine the clusters. This observation is the main insight behind Lloyd's algorithm.

8.3 Lloyd's Algorithm

Lloyd's algorithm, aka the K-means clustering algorithm, is perhaps the simplest unsupervised clustering method, which uses an *alternating* strategy that is very common in machine learning. The main idea is to (i) *pretend* that we know the centers and determine the clusters, (ii) *pretend* that we know the clusters and compute the centers, and then alternate between steps (i) and (ii) until convergence. More precisely, we start with some initial estimates $\hat{\mu}_1, \hat{\mu}_2, \dots, \hat{\mu}_K \in \mathbb{R}^D$, and then

- (i) Assign each datapoint to its closest center to produce a clustering:

$$\hat{C}_k = \left\{ i \in [N] : \|\hat{\mu}_k - \mathbf{x}_i\| \leq \|\hat{\mu}_\ell - \mathbf{x}_i\| \forall \ell \in [K] \right\}.$$

- (ii) Compute the center of each cluster:

$$\hat{\mu}_k = \frac{1}{|\hat{C}_k|} \sum_{i \in \hat{C}_k} \mathbf{x}_i.$$

Finally, we alternate between steps (i) and (ii) until convergence. See Figure 8.1 to build some intuition.

8.4 Initialization

As with most *alternating* algorithms, Lloyd's algorithm depends heavily on initialization, that is, the choice of initial centers $\{\hat{\mu}_k\}_{k=1}^K$. There are several popular options:

- **Random samples.** This option simply selects K data points randomly, that is, $\hat{\mu}_k = \mathbf{x}_i$ for some randomly chosen i . This tends to spread out initial centers.
- **Random partition.** This option first partitions data randomly into K clusters, and then computes the initial centers as the mean of each cluster. This tends to place all initial centers close to the center of the entire dataset.

- **K-means++**. This option aims to spread initial centers according to the data distribution. To this end, K-means++ selects one random data point \mathbf{x}_i as the first center $\hat{\boldsymbol{\mu}}_1$, and then for every $2 \leq k \leq K$, it chooses the k^{th} center from the remaining data with probability proportional to its closest existing center. That is, if \mathbf{x}_j is none of the first $k-1$ centers, then \mathbf{x}_j is chosen as the k^{th} center with probability proportional to

$$\min_{\ell \in [k-1]} \|\hat{\boldsymbol{\mu}}_\ell - \mathbf{x}_j\|^2.$$

With this initialization, Lloyd's algorithm is guaranteed to find a solution that is close (within a $\log K$ factor) to the optimal solution. This is remarkable, because the K-means problem is non-convex, NP-hard, so it is not evident that *any* algorithm should work.